

Computer Vision HW2 Report

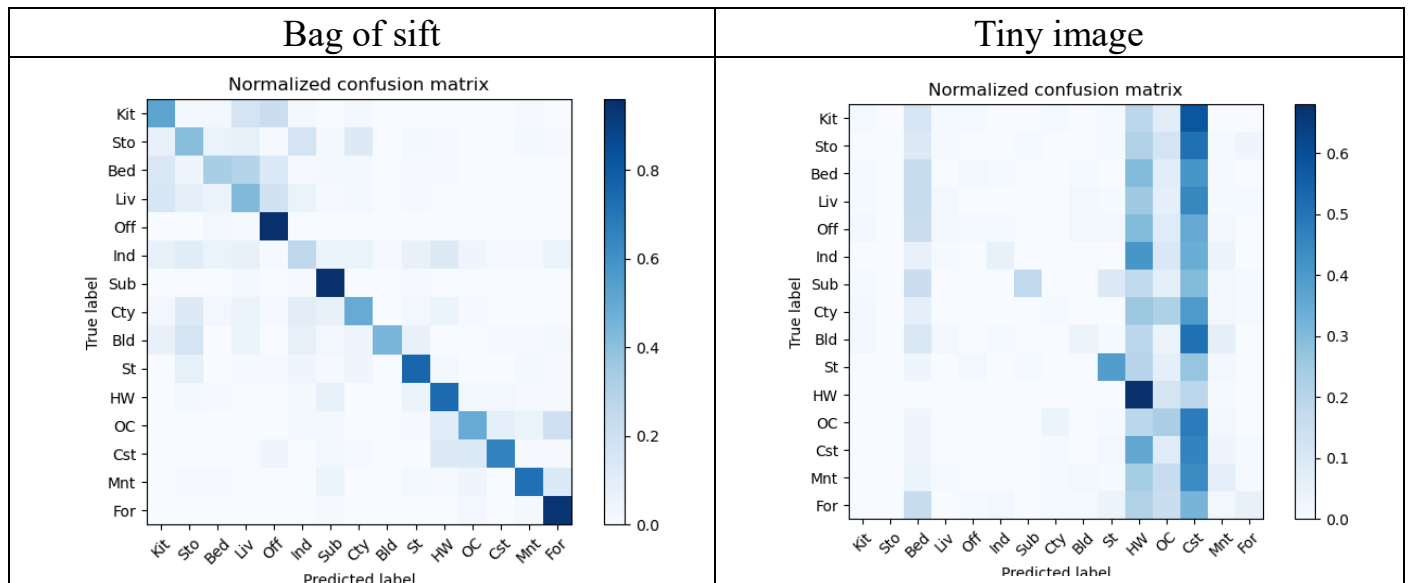
Student ID: R10921a36

Name: 石子仙

Part 1. (10%)

- Plot confusion matrix of two settings. (i.e. Bag of sift and tiny image representation) (5%)

Ans:



- Compare the results/accuracy of both settings and explain the result. (5%)

Ans:

- Setting of bag of sift and accuracy

	euclidean	cosine	jensenshannon	cityblock
k=5	0.523	0.555	0.5933	0.6073
k=7	0.533	0.56	0.6087	0.6073

- Setting of Tiny image and accuracy

	euclidean	cosine	jensenshannon	cityblock
k=5	0.157	0.157	0.142	0.188
k=7	0.162	0.158	0.145	0.189

上表標記紅色的為 best accuracy，綜合兩種表，我的程式碼最終設定為螢光筆底色的，k=7，cdist metric='cityblock'。

Confusion matrix 對角線上的點越深代表預測準確率越高，預測能力越好。可以從上表以及第一題的 confusion matrix 來看，Tiny image representation 的預測準確率比較

差，多半集中在 label HW 和 Cst。個人認為，有可能是僅從 tiny image 沒辦法提供足夠的資訊，而且也許只看一小張圖片的話，可能每張圖的 tiny image 都會有很相似的地方，一張圖片需要靠整張圖片的資訊合併起來才有意義，所以改把照片用 bag of sift 表示就能保留較多有用的資訊。

Part 2. (35%)

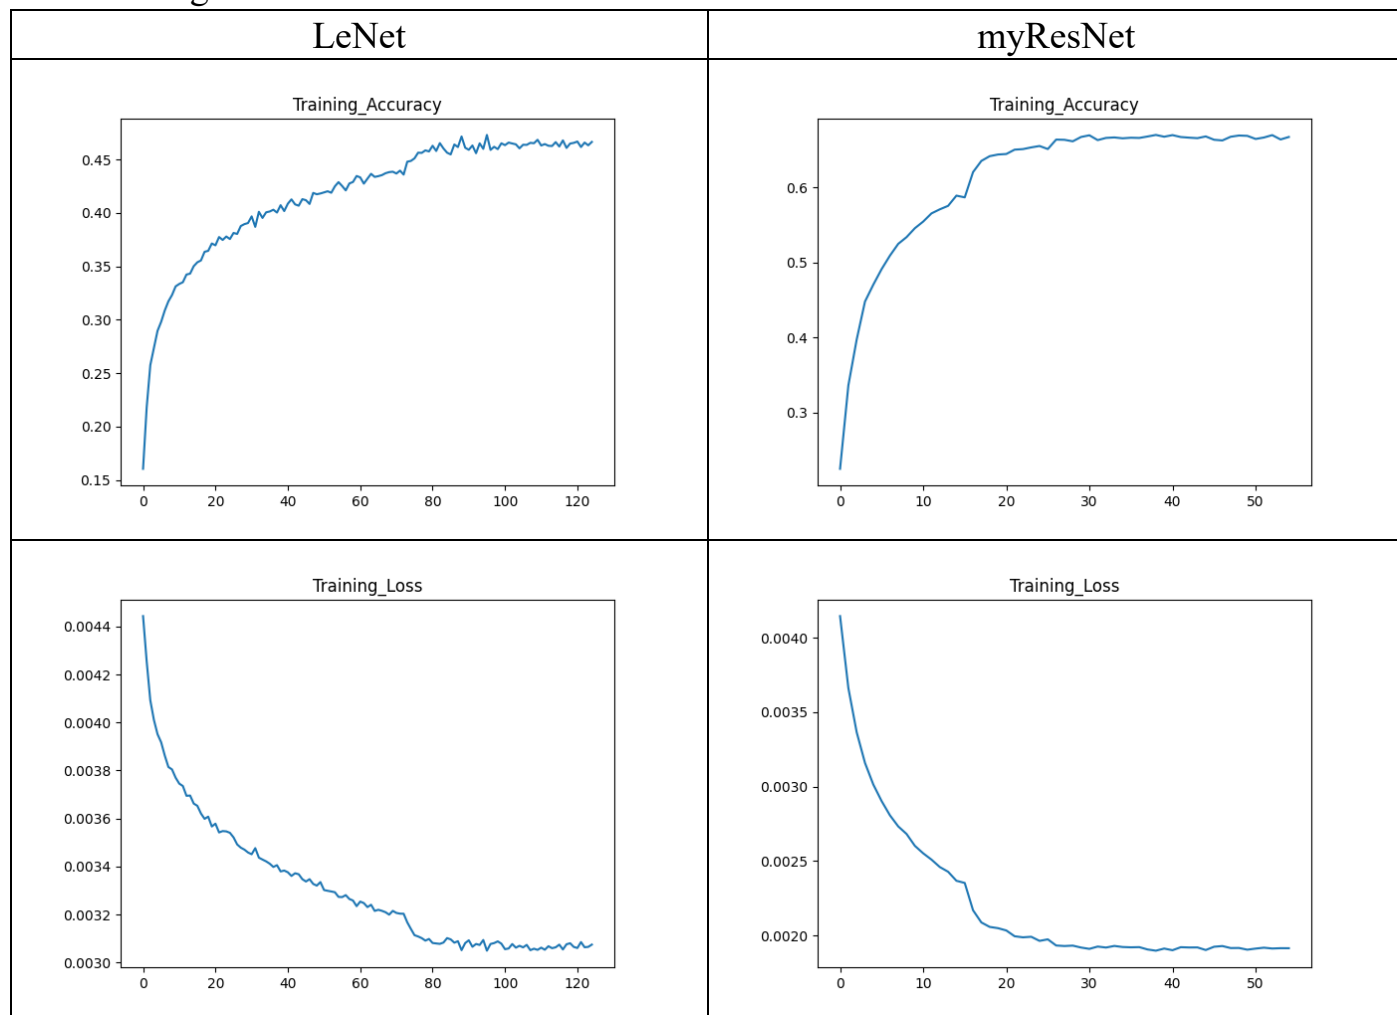
• Compare the performance on residual networks and LeNet. Plot the learning curve (loss and accuracy) on both training and validation sets for both 2 schemes. 8 plots in total. (20%)

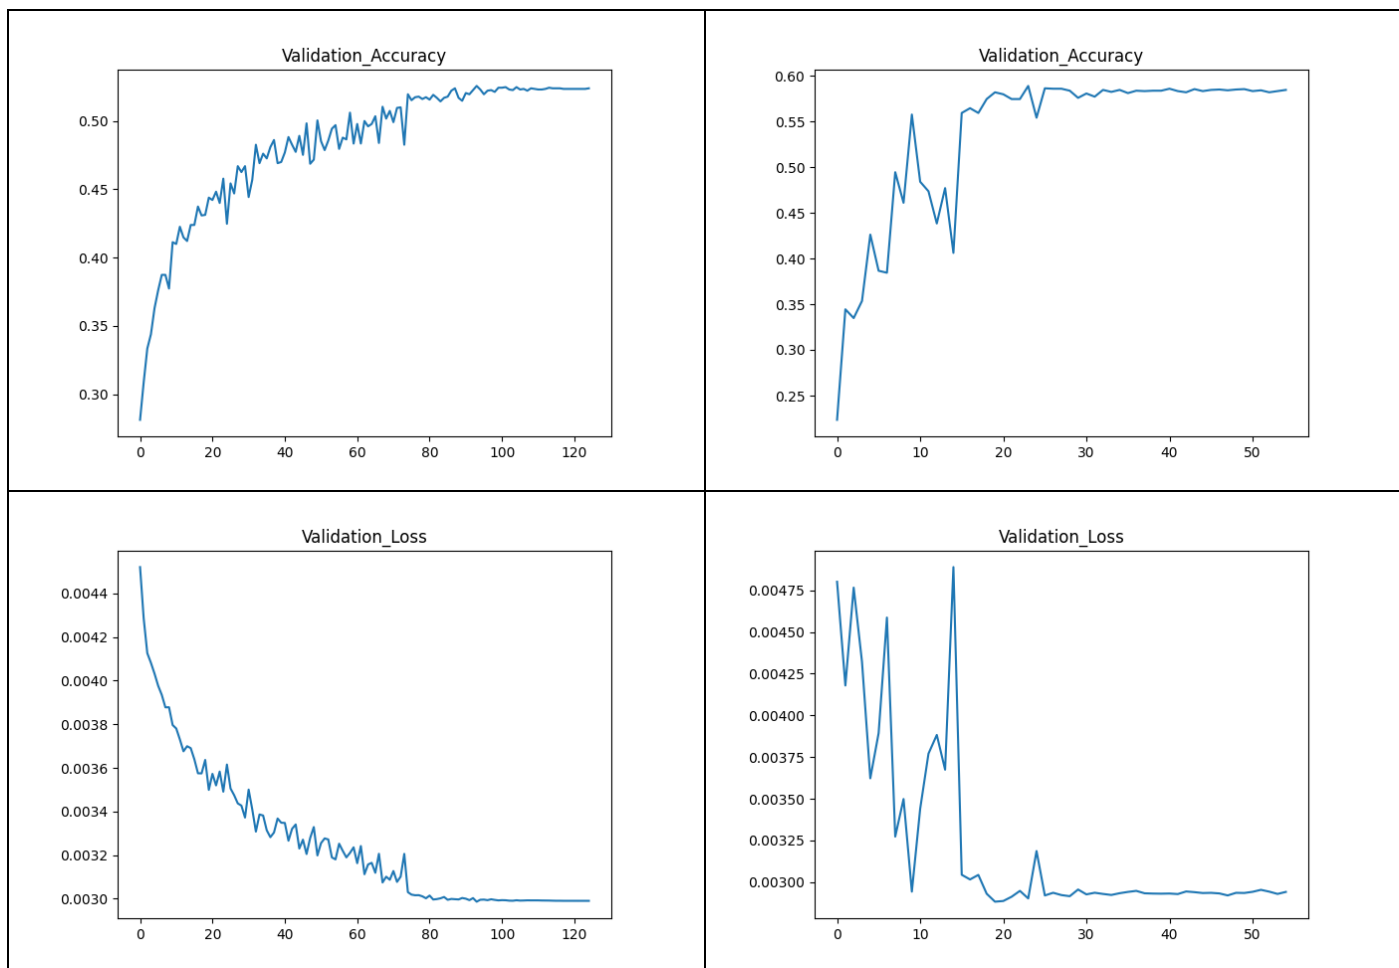
Ans:

● Accuracy

	LeNet	myResNet
Accuracy on public test	0.5992	0.6672

● Learning curve





在本題我有將原先助教提供的 optimizer 從 SGD 改成 Adam，scheduler 從 MultiStepLR 改成 ReduceLROnPlateau，若有連續 5 個 epoch model 沒有進步的話就調小 learning rate 為原本的 0.1 倍，並設定 early stop 為 30 個 epoch。


由上表 Accuracy 表格可得出 myResNet 的 performance 比 myLeNet 好很多(0.6672 v.s. 0.5992)，myResNet 的收斂速度比 myLenet 快。

• Attach basic information of the model you use including model architecture and number of the parameters. (5%)

Ans:

我使用 Pretrained resnet50，因為 model architecture 要佔好幾頁，這邊先用畫的，完整 print 出來的架構附在 report 最後。

● Model Architecture:



- **Number of the parameters:**

Total Trainable Params: 23,528,522

由於要把每一層都列出來會非常長，詳細附 report 的最後面。第一層 conv1 就是 $3*64*7*7 = 9408$ ，layer1.0.conv1.weight 是 $64*64 = 4096$ ，以此類推，最後再加總。

- **Briefly describe what method do you apply? (e.g. data augmentation, model architecture, loss function, semi-supervised etc.) (10%)**

Ans:

- 用來通過 baseline 的 model，我是使用 pretrained resnet50，並加入一些修改，最終 public test accuracy 為 0.8468。

- 我有做了以下改變：

1. 修改參數
2. 修改 optimizer 與 learning rate scheduler
3. 修改 data augmentation
4. 加入 semi-supervised

- 以下以各點詳細說明：

1. 修改參數

以下列的不是全部的 cfg，是比較重要的，另外我有加入 early_stop 來控制 training process。

```
'split_ratio': 0.9,  
'batch_size': 512,  
'num_epoch': 300,  
'early_stop': 30,  
'optimizer': 'Adam',  
'optim_hparas': {'lr': 0.001,  
                  'weight_decay': 0.001,  
                  'betas': (0.4, 0.999)}  
}
```

2. 修改 optimizer 與 learning rate scheduler

■ 將 Optimizer 從 SGD 改成 Adam，

```
'optimizer': 'Adam',  
'optim_hparas': {'lr': 0.001,  
                  'weight_decay': 0.001,  
                  'betas': (0.4, 0.999)}  
}
```

■ Learning rate scheduler 改成

```
optim.lr_scheduler.ReduceLROnPlateau(optimizer,  
                                     mode='min', patience=5, cooldown=3)
```

後面使用

```
scheduler.step(val_loss)
```



意思是如果 validation loss 在 5 個 epoch 下沒下降的話，就將 learning rate 改成原本的 0.1 倍（此為 default 的值我沒修改），改成這個後 accuracy 上升了約 0.01。



3. 修改 data augmentation

加入 torchvision.transforms 下的 AutoAugment、RandomRotation、RandomHorizontalFlip、RandomAffine

4. 使用 semi-supervised

在 validation accuracy 大於 0.7 且 validation accuracy 有進步時，就會開始使用 semi-supervised，令 confidence threshold 為 0.9，將 unlabeled data（這邊有將 public test 也拿來用）加入到 training set 中。由下圖的 learning curve 可以看到，在開始做 semi-supervised 後（validation accuracy 大於 0.7），training 與 validation 的 loss 與 accuracy 都有明顯的下降與提升。





Model Architecture 與 Number of parameters 的詳細

● Model Architecture:

詳細如下：

```
ResNet(
  (conv1): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3), bias=False)
  (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu): ReLU(inplace=True)
  (maxpool): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1, ceil_mode=False)
  (layer1): Sequential(
    (0): Bottleneck(
      (conv1): Conv2d(64, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
      (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (conv3): Conv2d(64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn3): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (relu): ReLU(inplace=True)
      (downsample): Sequential(
        (0): Conv2d(64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
```

```

    )
)
(1): Bottleneck(
  (conv1): Conv2d(256, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
  (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (conv3): Conv2d(64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn3): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (relu): ReLU(inplace=True)
)
(2): Bottleneck(
  (conv1): Conv2d(256, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
  (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (conv3): Conv2d(64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn3): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (relu): ReLU(inplace=True)
)
)
(layer2): Sequential(
  (0): Bottleneck(
    (conv1): Conv2d(256, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1),
bias=False)
    (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (conv3): Conv2d(128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)

```

```

        (bn3): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (relu): ReLU(inplace=True)
        (downsample): Sequential(
          (0): Conv2d(256, 512, kernel_size=(1, 1), stride=(2, 2), bias=False)
          (1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        )
      )
    (1): Bottleneck(
      (conv1): Conv2d(512, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
      (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (conv3): Conv2d(128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn3): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (relu): ReLU(inplace=True)
    )
    (2): Bottleneck(
      (conv1): Conv2d(512, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
      (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (conv3): Conv2d(128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn3): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (relu): ReLU(inplace=True)
    )
    (3): Bottleneck(
      (conv1): Conv2d(512, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

```



```

        (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
        (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (conv3): Conv2d(128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn3): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (relu): ReLU(inplace=True)
    )
)
(layer3): Sequential(
  (0): Bottleneck(
    (conv1): Conv2d(512, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1),
bias=False)
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu): ReLU(inplace=True)
    (downsample): Sequential(
      (0): Conv2d(512, 1024, kernel_size=(1, 1), stride=(2, 2), bias=False)
      (1): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
  )
  (1): Bottleneck(
    (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

```

```

        (relu): ReLU(inplace=True)
    )
    (2): Bottleneck(
      (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
      (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (relu): ReLU(inplace=True)
    )
    (3): Bottleneck(
      (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
      (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (relu): ReLU(inplace=True)
    )
    (4): Bottleneck(
      (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
      (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (relu): ReLU(inplace=True)
    )

```

```

)
(5): Bottleneck(
  (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
  (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (relu): ReLU(inplace=True)
)
)
(layer4): Sequential(
  (0): Bottleneck(
    (conv1): Conv2d(1024, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1),
bias=False)
    (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (conv3): Conv2d(512, 2048, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(2048, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu): ReLU(inplace=True)
    (downsample): Sequential(
      (0): Conv2d(1024, 2048, kernel_size=(1, 1), stride=(2, 2), bias=False)
      (1): BatchNorm2d(2048, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
  )
  (1): Bottleneck(
    (conv1): Conv2d(2048, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)

```

```

        (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (conv3): Conv2d(512, 2048, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn3): BatchNorm2d(2048, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (relu): ReLU(inplace=True)
    )
    (2): Bottleneck(
        (conv1): Conv2d(2048, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
        (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (conv3): Conv2d(512, 2048, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn3): BatchNorm2d(2048, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (relu): ReLU(inplace=True)
    )
)
(avgpool): AdaptiveAvgPool2d(output_size=(1, 1))
(fc): Sequential(
  (0): Linear(in_features=2048, out_features=10, bias=True)
)
)

```

● 詳細 number of model parameter :

Modules	Parameters
conv1.weight	9408
bn1.weight	64
bn1.bias	64
layer1.0.conv1.weight	4096
layer1.0.bn1.weight	64
layer1.0.bn1.bias	64

layer1.0.conv2.weight	36864	
layer1.0.bn2.weight	64	
layer1.0.bn2.bias	64	
layer1.0.conv3.weight	16384	
layer1.0.bn3.weight	256	
layer1.0.bn3.bias	256	
layer1.0.downsample.0.weight	16384	
layer1.0.downsample.1.weight	256	
layer1.0.downsample.1.bias	256	
layer1.1.conv1.weight	16384	
layer1.1.bn1.weight	64	
layer1.1.bn1.bias	64	
layer1.1.conv2.weight	36864	
layer1.1.bn2.weight	64	
layer1.1.bn2.bias	64	
layer1.1.conv3.weight	16384	
layer1.1.bn3.weight	256	
layer1.1.bn3.bias	256	
layer1.2.conv1.weight	16384	
layer1.2.bn1.weight	64	
layer1.2.bn1.bias	64	
layer1.2.conv2.weight	36864	
layer1.2.bn2.weight	64	
layer1.2.bn2.bias	64	
layer1.2.conv3.weight	16384	
layer1.2.bn3.weight	256	
layer1.2.bn3.bias	256	
layer2.0.conv1.weight	32768	
layer2.0.bn1.weight	128	
layer2.0.bn1.bias	128	
layer2.0.conv2.weight	147456	
layer2.0.bn2.weight	128	
layer2.0.bn2.bias	128	
layer2.0.conv3.weight	65536	
layer2.0.bn3.weight	512	
layer2.0.bn3.bias	512	
layer2.0.downsample.0.weight	131072	
layer2.0.downsample.1.weight	512	
layer2.0.downsample.1.bias	512	
layer2.1.conv1.weight	65536	

layer2.1.bn1.weight	128	
layer2.1.bn1.bias	128	
layer2.1.conv2.weight	147456	
layer2.1.bn2.weight	128	
layer2.1.bn2.bias	128	
layer2.1.conv3.weight	65536	
layer2.1.bn3.weight	512	
layer2.1.bn3.bias	512	
layer2.2.conv1.weight	65536	
layer2.2.bn1.weight	128	
layer2.2.bn1.bias	128	
layer2.2.conv2.weight	147456	
layer2.2.bn2.weight	128	
layer2.2.bn2.bias	128	
layer2.2.conv3.weight	65536	
layer2.2.bn3.weight	512	
layer2.2.bn3.bias	512	
layer2.3.conv1.weight	65536	
layer2.3.bn1.weight	128	
layer2.3.bn1.bias	128	
layer2.3.conv2.weight	147456	
layer2.3.bn2.weight	128	
layer2.3.bn2.bias	128	
layer2.3.conv3.weight	65536	
layer2.3.bn3.weight	512	
layer2.3.bn3.bias	512	
layer3.0.conv1.weight	131072	
layer3.0.bn1.weight	256	
layer3.0.bn1.bias	256	
layer3.0.conv2.weight	589824	
layer3.0.bn2.weight	256	
layer3.0.bn2.bias	256	
layer3.0.conv3.weight	262144	
layer3.0.bn3.weight	1024	
layer3.0.bn3.bias	1024	
layer3.0.downsample.0.weight	524288	
layer3.0.downsample.1.weight	1024	
layer3.0.downsample.1.bias	1024	
layer3.1.conv1.weight	262144	
layer3.1.bn1.weight	256	

layer3.1.bn1.bias	256	
layer3.1.conv2.weight	589824	
layer3.1.bn2.weight	256	
layer3.1.bn2.bias	256	
layer3.1.conv3.weight	262144	
layer3.1.bn3.weight	1024	
layer3.1.bn3.bias	1024	
layer3.2.conv1.weight	262144	
layer3.2.bn1.weight	256	
layer3.2.bn1.bias	256	
layer3.2.conv2.weight	589824	
layer3.2.bn2.weight	256	
layer3.2.bn2.bias	256	
layer3.2.conv3.weight	262144	
layer3.2.bn3.weight	1024	
layer3.2.bn3.bias	1024	
layer3.3.conv1.weight	262144	
layer3.3.bn1.weight	256	
layer3.3.bn1.bias	256	
layer3.3.conv2.weight	589824	
layer3.3.bn2.weight	256	
layer3.3.bn2.bias	256	
layer3.3.conv3.weight	262144	
layer3.3.bn3.weight	1024	
layer3.3.bn3.bias	1024	
layer3.4.conv1.weight	262144	
layer3.4.bn1.weight	256	
layer3.4.bn1.bias	256	
layer3.4.conv2.weight	589824	
layer3.4.bn2.weight	256	
layer3.4.bn2.bias	256	
layer3.4.conv3.weight	262144	
layer3.4.bn3.weight	1024	
layer3.4.bn3.bias	1024	
layer3.5.conv1.weight	262144	
layer3.5.bn1.weight	256	
layer3.5.bn1.bias	256	
layer3.5.conv2.weight	589824	
layer3.5.bn2.weight	256	
layer3.5.bn2.bias	256	

layer3.5.conv3.weight	262144	
layer3.5.bn3.weight	1024	
layer3.5.bn3.bias	1024	
layer4.0.conv1.weight	524288	
layer4.0.bn1.weight	512	
layer4.0.bn1.bias	512	
layer4.0.conv2.weight	2359296	
layer4.0.bn2.weight	512	
layer4.0.bn2.bias	512	
layer4.0.conv3.weight	1048576	
layer4.0.bn3.weight	2048	
layer4.0.bn3.bias	2048	
layer4.0.downsample.0.weight	2097152	
layer4.0.downsample.1.weight	2048	
layer4.0.downsample.1.bias	2048	
layer4.1.conv1.weight	1048576	
layer4.1.bn1.weight	512	
layer4.1.bn1.bias	512	
layer4.1.conv2.weight	2359296	
layer4.1.bn2.weight	512	
layer4.1.bn2.bias	512	
layer4.1.conv3.weight	1048576	
layer4.1.bn3.weight	2048	
layer4.1.bn3.bias	2048	
layer4.2.conv1.weight	1048576	
layer4.2.bn1.weight	512	
layer4.2.bn1.bias	512	
layer4.2.conv2.weight	2359296	
layer4.2.bn2.weight	512	
layer4.2.bn2.bias	512	
layer4.2.conv3.weight	1048576	
layer4.2.bn3.weight	2048	
layer4.2.bn3.bias	2048	
fc.0.weight	20480	
fc.0.bias	10	

+-----+-----+

Total Trainable Params: 23528522